

CHIROVOX

Thanks for the nudge—I'll give you both: a quick explanation of why we use `Dense(128, activation='relu')` and then a full, step-by-step breakdown of your CHIROVOX project: architecture, algorithms, libraries, training flow, and how each file fits in.

Why we use `keras.layers.Dense(128, activation='relu')`

- What it is: A fully connected layer with 128 neurons and ReLU nonlinearity.
- Where it sits: Right after `Flatten()`, just before the final softmax classifier in [train.py](#).
- What it does:
 - Learns high-level, non-linear combinations of the features extracted by the Conv2D/MaxPool layers.
 - ReLU introduces non-linearity to model complex decision boundaries.
 - 128 units is a good balance for a small model (not too tiny, not too heavy) for 64×64 inputs.
- Effect on training:
 - Acts as a “feature combiner” before the final classification.
 - Helps the softmax layer see richer, more separable features.
- Rough math: For input feature vector x (from `Flatten`), `Dense` computes $y = \text{ReLU}(Wx + b)$ with $W \in \mathbb{R}[F \times 128]$, then the final `Dense(num_classes, softmax)` maps y to class probabilities.

Tip: If you see overfitting later, consider adding `Dropout` (e.g., `Dropout(0.3)`) before/after this `Dense` layer.

CHIROVOX: Full project overview

What it is

A sign-language helper that:

- Uses a CNN trained on gesture images
- Converts the model to `TensorFlow.js`
- Runs live camera detection in the browser

- Speaks the recognized sign in English/Sinhala/Tamil
- Includes learning (education) and dictionary pages

High-level architecture

- Training (Python, offline):
 - [train.py](#) trains a small CNN on images under pages/gesture/dataset/...
 - Exports a Keras model ([gesture_model.h5](#)) and [labels.json](#)
 - Converts to TFJS (writes [model.json](#) + shard files) under [model](#)
- Inference (browser, online/offline):
 - [gesture.js](#) loads the TFJS model
 - Captures camera frames (MediaPipe camera utils + <video>)
 - Optionally draws hands (MediaPipe Hands) and overlays
 - Resizes frames to 64×64 and feeds to the TFJS model
 - Speaks the predicted label using Web Speech Synthesis
 - Falls back to a finger-count heuristic if the model isn't available
- UI/UX:
 - [index.html](#) landing with header, language selector, theme toggle
 - [education.html](#) learning cards
 - [dictionary.html](#) table of common phrases
 - Shared theme in [theme.css](#), plus page CSS
 - Global language state via [lang.js](#) synced across pages

Algorithms used

1) Image classification (CNN)

Defined in [train.py](#):

- Input: 64×64 RGB images of hands (from folder structure)
- Data pipeline: [ImageDataGenerator](#) with rescaling and light augmentation
- Architecture:
 - Conv2D(32) + MaxPool
 - Conv2D(64) + MaxPool

- Conv2D(128) + MaxPool
- Flatten
- Dense(128, relu) ← feature combiner
- Dense(num_classes, softmax)
- Loss: categorical_crossentropy
- Optimizer: Adam
- Epochs: 10 (can increase if not overfitting)

2) TFJS conversion

- Uses tensorflowjs_converter to export TFJS artifacts for the web app.
- Client loads [model.json](#) and predicts per-frame.

3) Live detection pipeline

In [gesture.js](#):

- Camera frames → canvas → resize to 64×64 → tf.browser.fromPixels → model.predict → argmax → label
- Optional fallback when TFJS model not found: heuristic finger counting via MediaPipe landmarks (maps finger count to a small set of gestures).
- TTS: Map predicted label to localized text and speak via SpeechSynthesis with language-specific voices.

Libraries used

Python (training)

- tensorflow==2.15
- tensorflowjs (for conversion)
- numpy
- matplotlib (optional plotting if added later)
- Standard library: os, json, subprocess, sys, pathlib

Browser (inference/UI)

- TensorFlow.js (CDN): @tensorflow/tfjs
- MediaPipe: [hands](#), camera_utils, drawing_utils (CDN)
- Web Speech API: [window.speechSynthesis](#) (no extra lib)

- Vanilla JS/HTML/CSS

Project structure mapped to behavior

- [train.py](#)
 - Loads images from pages/gesture/dataset/{class_name}/...jpg
 - Augments, builds, trains CNN
 - Saves [gesture_model.h5](#)
 - Writes [labels.json](#) containing ordered class names
 - Converts model to TFJS into [model](#) (so browser can load)
- [requirements.txt](#)
 - Tooling to set up Python training environment
- [gesture.html](#)
 - Loads theme, TFJS, MediaPipe scripts
 - Displays orb, “Open” button, language select
 - Has <video id="camera"> and <canvas id="output"> for previews/overlays
- [gesture.js](#)
 - Model loading logic (tries multiple paths for [model.json](#))
 - Inference on frames (64×64)
 - Rendering of landmarks (MediaPipe) and overlay drawing
 - Fallback heuristic if no model
 - Throttled TTS, language mapping, event sync
 - Theme/localStorage integration
- [gesture.css](#)
 - Camera area styling, buttons, layout, responsive tweaks
- [lang.js](#)
 - Global language store (localStorage)
 - Emits chirovox:langchange events so pages stay in sync
- [index.html](#)
 - Header with branding, language selector, theme toggle

- Calls [lang.js](#) and wires theme/light-mode
 - Quick links to gesture/education/dictionary
- [education.html](#)
 - Title (“Learn Sri Lankan Sign Language”)
 - Language selector (synced)
 - Cards (image + heading + description); we enlarged images recently
- [dictionary.html](#)
 - Language selector (synced)
 - Glass-styled table of phrases
- [theme.css](#) and [style.css](#)
 - Global theme and page styles
 - Light/dark backgrounds (CHIROVOXWHITE.png for light mode)
 - Premium controls (gold language selector), layout polish
 - Education card grid; we updated image sizes
- [quiz.html](#), [schools.html](#), [about.html](#)
 - Additional content pages with shared theme and structure
 - Schools page includes a quick speak section
 - About includes project description and credits

Training: step-by-step how you did it (and how to re-run)

1. Prepare dataset
 - Create class folders:
 - [hello](#), iloveyou, yes, no, etc.
 - Put sample images in each folder (JPG/PNG). More balanced data = better results.
2. Create/activate Python environment
 - On Windows PowerShell:
 -
 -

-
-

3. Run training

- From project root:

-
-
-
-

- What it does:

- Loads images with augmentation
- Trains the CNN
- Saves gesture_model.h5
- Writes labels.json to labels.json
- Converts to TFJS and writes model.json + shards to model

4. Start local server for testing

- From project root:

-
-
-
-

- Visit in browser:

- Laptop: <http://localhost:8000/pages/gesture/gesture.html>
- Phone (same Wi-Fi): http://<your-laptop-ip>:8000/pages/gesture/gesture.html
 - You may need HTTPS (use ngrok) for camera permissions on some devices

5. Use the camera page

- Click Open
- Show a known gesture

- You should see the class name and hear TTS in your chosen language

How predictions flow in the browser (detailed)

- [gesture.js](#):
 - On Open:
 - Gets camera stream (MediaPipe Camera helper or getUserMedia fallback)
 - For each frame:
 - Draws the frame to canvas for preview
 - Converts frame to tensor:
 - `tf.browser.fromPixels` → resize to [64,64] → [0..1] normalization → [1, H, W, 3]
 - Runs [model.predict\(t\)](#) → argMax → predicted label
 - Displays label → maps to localized string → speaks with correct language voice (SpeechSynthesis)
 - If TFJS model fails to load:
 - Loads labels.json if available (for UI)
 - Uses a simple finger count heuristic to pick a plausible label

Improvements you can consider (optional)

- Add Dropout(0.3) before/after Dense(128) to reduce overfitting
- Add EarlyStopping and ModelCheckpoint callbacks
- Train longer (20–30 epochs) and monitor val accuracy
- Use transfer learning (e.g., MobileNetV2) for higher accuracy
- Use MediaPipe landmarks as input features for a small MLP (fast, low data)
- More consistent dataset lighting and orientations
- Balance class samples to avoid bias

Quick FAQ

- Why 64×64 input?
 - Keeps the model small for fast TFJS inference on phones. Increase if accuracy requires.

- Why we used ImageDataGenerator?
 - Simple augmentation improves generalization without adding code complexity.
- Why MediaPipe if we classify raw images?
 - MediaPipe helps visualize hand skeleton and provides a fallback heuristic. You can also switch to using landmark vectors as inputs.
- Why is HTTPS recommended on phone?
 - Browser camera APIs often require secure context (HTTPS) unless using localhost or some private LAN exceptions.